

# Plug in the Safety Chip: Enforcing Constraints for LLM-driven Robot Agents

Ziyi Yang, Shreyas S. Raman, Ankit Shah<sup>†</sup>, and Stefanie Tellex  
Department of Computer Science, Brown University, United States

**Abstract**—Recent advancements in large language models (LLMs) have enabled a new research domain, LLM agents, for solving robotics and planning tasks by leveraging the world knowledge and general reasoning abilities of LLMs obtained during pretraining. However, while considerable effort has been made to teach the robot the “dos”, the “don’ts” received relatively less attention. We argue that, for any practical usage, it is as crucial to teach the robot the “don’ts”: conveying explicit instructions about prohibited actions, assessing the robot’s comprehension of these restrictions, and, most importantly, ensuring compliance. Moreover, verifiable safe operation is essential for deployments that satisfy worldwide standards such as ISO 61508, which defines standards for safely deploying robots in industrial factory environments worldwide. Aiming at deploying the LLM agents in a collaborative environment, we propose a queryable safety constraint module based on linear temporal logic (LTL) that simultaneously enables natural language (NL) to temporal constraints encoding, safety violation reasoning and explaining, and unsafe action pruning. To demonstrate the effectiveness of our system, we conducted experiments in VirtualHome environment and on a real robot. The experimental results show that our system strictly adheres to the safety constraints and scales well with complex safety constraints, highlighting its potential for practical utility.

## I. INTRODUCTION

Recently, Large Language Models (LLMs) have enabled LLM-based autonomous agents, or more recently termed LLM agents [1], to emerge as a promising approach for various applications [2], including planning, logical reasoning, as well as robotic tasks. However, with the increasing desire to deploy these agents in **daily settings** for robotic tasks, ensuring safety has become an inevitable concern, particularly in situations where safety holds more significance than the assigned tasks themselves. Safety concerns are particularly important in industrial or factory settings, where strict safety standards such as ISO 61508 [3] must be met and maintained for automation systems to be deployed. In this work, inspired by the categorization of Gu et al. [4], we consider an LLM agent to be safe if it acts, reasons, and generalizes obeying human desires and never reaches unsafe states. Consider a scenario in which a physically embodied robot agent is being deployed across various environments for different roles, such as housekeeping for the elderly in a nursing home or drug delivery for patients in a hospital. While current planning and control mechanisms are mostly interested in the robot’s capabilities, there are vital concerns underlying safety issues in these specific domains. To answer

their question, “**is this robot safe?**” an ideal solution would be a “**safety chip**” that can be plugged into existing robot agents: this chip would empower the robot agent with the skills to comprehend customized safety instructions per user requests, convey safety specifications in its belief through NL, and adhere to established safety standards. Still, three major obstacles exist on the path toward this “safety chip” for LLM-based robot agents:

- The inherently **probabilistic nature** of LLM agents hinders their ability to consistently adhere to safety standards. This problem becomes exacerbated when goal specifications are specified by untrained users at runtime in natural language, or conflict with safety constraints.
- LLMs struggle to scale up as the complexity of constraints increases, which can also distract an LLM agent from completing the original task.
- Currently, LLM agents rely on external modules such as affordance model for grounded decision-making. Despite their high in-domain performance, such pretrained modules are likely to have limited ability to generalize to new domains or be customized to human preference.

To overcome the aforementioned challenges, our work introduces a hybrid system where safety constraints are represented in terms of a formal language, Linear Temporal Logic (LTL), which can be verified to be correct. We map between English instructions describing the safety constraints and LTL constraints, which are then verified by experts, relying on the fact that verifying LTL expressions is easier and more reliable than writing them from scratch [5]. This approach provides a pathway where the expression of the safety constraints could in principal meet ISO 61508 standards [3] although actually meeting those standards is beyond the scope of this work. Our overall approach can operate with any existing language understanding framework regardless of its technical underpinnings to provide safety guarantees, as demonstrated in Figure 1. In addition to this “black box” safety guarantee, our approach, when combined with an LLM framework, can use reprompting to find new plans to achieve task success without violating safety constraints. The main contributions of this work are:

- Proposed a safety constraint module for customizable constraints and integrated the proposed module into an existing LLM agent framework.
- A fully prompting-based approach that supports predicate syntax for translating NL to LTL and explaining violation of LTL specification in NL.

Code and appendix are available at <https://yzylmc.github.io/safety-chip/>

<sup>†</sup> Author was at Brown University when the work was done.

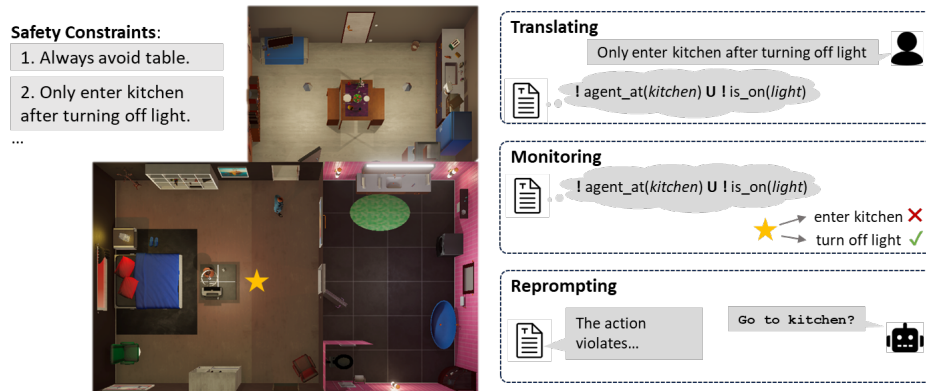


Fig. 1: Demonstration of safety chip’s functionality of translating NL to LTL formulas, monitoring the agent’s decision-making process, and reprompting with reason of violation for re-planning.

- A formal method-based action pruning and feedback for active re-planning for LLM agents.
- Deployed the system in an embodied environment and on real robot platforms for baseline comparisons.

## II. RELATED WORKS

### A. LTL for Safety in Robotics

Linear temporal logic (LTL) [6] has found utility in expressing temporal task specification for various planning and learning tasks [7, 8, 9, 10]. Its application for safety purposes arises from its expressivity and unambiguous semantics. Runtime verification and monitoring for pre-defined safety constraints specified as temporal logic formulas has been studied extensively [11, 12, 13, 14, 15], and runtime monitoring along with specification elicitation was recognized as an important challenge in formalizing human-robot interaction by Kress-Gazit et al. [16]. Our approach to safety violation identification and constraint enforcement is inspired by runtime monitoring frameworks reliant on temporal logic specifications. Mao et al. [17] demonstrated using a variant of LTL for safety in programmable logic controllers (PLCs) commonly used for robotic control in a way consistent with ISO 61508, showing its potential for use in the context of industrial robotic safety.

In the broader context of Human-Robot Interaction (HRI), the subject of safety has been investigated in-depth [18]. Here safety is viewed through a shared physical workspace, we view safety through task planning as language-specified constraints that have to be satisfied. To the best of our knowledge, such task-level safety satisfaction is under-explored by existing HRI research.

### B. LLM Agents

LLMs exhibit substantial reasoning abilities and have been deployed for planning in either embodied domain or on real robots, and a line of works [19, 20, 21, 22, 23] have shown interesting results on deploying the LLM agents for robotics. Among them, Huang et al. [24] has touched on the safety topic by providing a neural network-based safety module in parallel with affordance and human preference modules together as scoring functions, and Singh et al. [25], Liang

et al. [26] proposed the use of python code to express robot plans. Recent works [27, 28] turn to Planning Domain Definition Language (PDDL) for more accurate reasoning and solving longer-horizon planning problems, while a big assumption they made is that PDDL’s syntax is sufficiently expressive for various NL instructions.

Our method attempts to combine the strength of both paradigms by plugging the proposed safety module into an existing LLM agent that utilizes natural language as an interface. We hope to adapt the generality and expressiveness of natural language and the rigorousness of formal language.

### C. Translation between NL and LTL

The attempts to translate NL to LTL range from traditional RNN model [29, 30] to the recent LLM-based works [31, 32, 33]. However, a shared challenge in these studies is the limited availability of training data. While LLMs show great improvement in translation, their performance deteriorates with increased complexity, posing difficulties in generalizing to out-of-distribution formulas.

LTL to NL works such as Cherukuri et al. [34] are still in an early stage, likely due to the challenge of learning representations of high-level abstractions of LTL and automata. Our approach tackles the problem by providing state information as a knowledge representation of the automaton. We believe such backward translation could serve as a valuable interface between humans and LTL-based robot systems.

## III. METHOD

### A. Problem Definition

Our aim is to create a language understanding system capable of meeting strict safety standards that might be applied in an industrial setting. We assume access to an existing language understanding system, such as an LLM agent, that maps between language and robot action using any mechanism, but that system has no constraints on using a particular formal representation or safety guarantees. We additionally assume access to a human designer who can express safety constraints in English, and has the capability to verify those constraints once expressed in a formal language.

## B. Overview

Our proposed approach is to create a “safety chip” that can take a language understanding system that does not have safety guarantees and provide a framework that can provide verifiable guarantees that safety constraints will not be violated. Moreover, we provide a reprompting strategy to enable LLM-based agents in particular, to plan to avoid violating the safety constraints. Our “safety chip” consists of a human-agent team. The human specifies safety constraints in English. The robot then maps those constraints to a formal representation. The representation is then verified by the human for correctness. Once correctness has been verified, we enforce these constraints on the existing LLM Agent. Specifically, for an existing LLM agent making decisions based on in-context knowledge  $l_{agent}$ , we enforce a series of constraints in natural language  $\{c_i\}$  concerning safety aspects by encoding each of them into LTL formulas  $\{\varphi_i\}$  and composing them together into one single formula  $\Phi = \bigwedge_{i=0}^{\infty} \varphi_i$ , which is then stored as an automaton,  $A$ . To monitor the LLM agent executing a high-level task, which ought to be decomposed by the agent into a set of subtasks or actions  $\{t_i\}$ , we mask unsafe action  $t_i$  at time step  $i$  by progressing its partial trajectory  $T_i = \{t_0, t_1, \dots, t_i\}$  with  $A$ , and generate a natural language query  $Q$  regarding the constraint violation induced by the unsafe action. Responding to the safety constraint query, our query system generates an explanation  $l_A$  from  $P(l_A|Q, c_{1:i}, T_i)$ . To regenerate a safe action  $t'_i$ , we prompt the LLM agent again with the explanation  $l_A$  to modify its distribution from  $P(t'_i|l_{agent}, T_{i-1})$  to  $P(t'_i|l_{agent}, T_{i-1}, l_A)$ .

As Figure 2 shows, our system consists of a translator system for NL to LTL translation, a query system for explaining safety violations, and a constraint monitor system for validating and pruning actions generated by the LLM agent. The automaton serves as the central part of the safety constraints module: it represents the safety constraints encoded from the LTL formula in a validatable form, reasons the violation of constraints with state changes of propositions, and validates the agent’s proposed action by progressing the proposition-level trajectory. In addition, the output of the query system is also used to assist the agent in performing re-planning.

## C. Identifying Safety Constraints

We propose a two-phase approach for identifying and verifying safety constraints. First, we asked the human designer to specify the constraints  $\{c_i\}$  using natural language. We automatically map the language to LTL formulas  $\{\varphi_i\}$ . Then we ask them to verify the formula is correct. We adopted the modular framework of Lang2LTL [33] as NL to domain-specific LTL formula **translator** using a predefined vocabulary. This involves extracting referring expressions, grounding referring expressions to propositions within the vocabulary, translating lifted utterances to formulas, and finally producing grounded formulas by replacing placeholders with grounded propositions (see Figure 3). The noteworthy distinction is that our translation module relies solely on in-context learning and necessitates no fine-tuning due to the compositionality nature of our approach: the safety

constraints are assumed to be provided as a series of basic segments. This enables separate translation and assembly into a single LTL formula  $\Phi$  with logical operators, which is usually logical AND (&) since the constraints must be satisfied simultaneously. Besides, we also have support for predicates that can essentially enable the translation system to go beyond the navigation domain.

For verifying the formula is correct, we rely on the human having knowledge of the specific formal language being used. This is required in order to meet the strict safety standards specified by ISO 61508 [3]. Greenman et al. [5] show that even for experts it is easier to verify a formula than it is to write a new formula from scratch. Furthermore, in our future work, we plan to explore integrated dialog systems that can enable untrained users to specify safety constraints using language and question-asking dialog to meet very high levels of certainty that they are correctly understood.

## D. Enforcing Safety Constraints

Pruning unsafe actions performed by the **constraint monitor** module constitutes the core functionality of the safety chip for enforcing safety constraints. For a sequence of actions, we keep tracking the truth value of each proposition at and between each action, and we determine the validity of the action sequence by progressing it through an automaton  $A$ : we consider a trajectory,  $T$ , to be valid if all state changes are valid and the final state of the sequence has a path to the accepting states of the automaton, or if the final state falls in the accepting states, if the agent terminates the execution as its last action.

Our approach for runtime monitoring with LTL is compatible with any language understanding system, regardless of its internal structure capabilities of the language understanding system, the safety chip will simply stop the robot’s execution if it detects it is about to move to an unsafe state. On the other extreme, if the internal system also uses LTL, we can follow the approach of Raman et al. [35] and give the user detailed feedback about what constraints are violated as soon as the runtime language has been interpreted, for example by a system such as Lang2LTL [33]. However, LTL lacks expressivity as compared to natural language, and systems may wish to incorporate the power of large language models without being limited by the expressivity of LTL. For such hybrid systems, we present a strategy to reprompt the LLM using information from the safety chip, described in the following two sections.

## E. Reprompting with Large Language Models

Since we keep tracking the truth values of each proposition in the LTL formula as the LLM agent is planning and executing, state transition in the automaton can also be monitored. Thus, assuming we have perfect truth value functions and transition model, any violation of the safety constraints can be detected in advance, and we can mask out the unsafe action  $t_i$  and ask the LLM agent to re-generate the plan from the last step. However, masking alone doesn’t provide any information to the LLM agent, causing the agent to likely

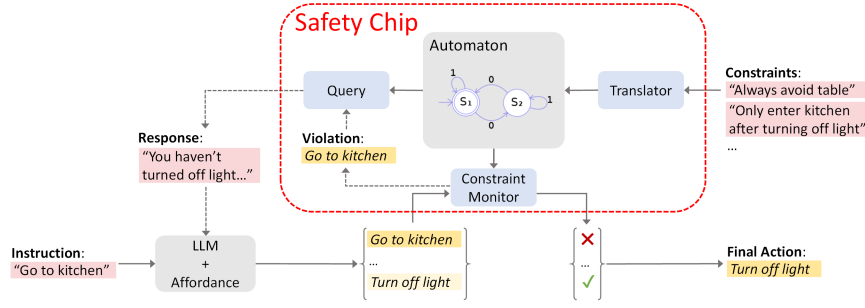


Fig. 2: System Diagram. The safety chip infers specification of constraints from NL and stores it as automaton. The automaton is then used for monitoring the decision-making process of an LLM agent, and responses to agent queries can be generated accordingly to help the LLM agent with reprompting.

choose an action within the same distribution of the masked unsafe action since its distribution remains unchanged. In order to shift the distribution towards the ground truth or safe actions, we want to provide information through the context of the prompt, so that the agent can make a decision based on the updated context, and we term this process *reprompting*.

An example implementation of reprompting in the decision-making loop is shown as follows (text in grey is the prefix in the prompt, black is generated by the LLM agent, and green is explanation  $l_A$  generated for reprompting. task specification and examples are removed for simplicity):

Go to toilet  
 Description: Travel through the house and locate the bathroom, and proceed to the toilet.  
 0. walk to bathroom  
 Error: The user is trying to enter the bathroom before entering the living room, which violates the constraint “you have to enter living room before bathroom”. The correct plan would be:  
 0. walk to living room  
 Error: The action “walk to livingroom” violates the constraint “you have to enter bedroom before going into living room”. According to the state change, the user is currently in the kitchen and has not entered the bedroom yet. The correct plan would be:  
 0. walk to bedroom  
 1. walk to living room  
 2. walk to bathroom  
 3. walk to toilet  
 4. DONE

#### F. Queryable Safety Constraints

The **query system** serves the purpose of assisting LLM agents in effective reprompting. Inspired by [36, 37], in addition to masking out unsafe actions, we utilize the query system to provide closed-loop feedback to the agent regarding reasons for violation. This approach prompts the query system to reason over explanations for validation results and feeds the output to the LLM agent, introducing an inductive bias from the query system. Similar to the translation system, the query system relies entirely on prompting. As recent

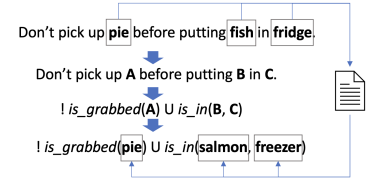


Fig. 3: NL to LTL translation example. Referring expressions are extracted and grounded into predefined propositions, and lifted NL utterances are translated into a lifted formula, then combined to produce the final result.

research [38] shows, LLMs are susceptible to hallucination and generally unreliable for classification tasks, making them unsuitable to be directly applied for validation detection in our case. Fortunately, the automaton can provide action validity through trajectory progression, thereby reducing LLMs’ role to explanation. For that, we construct agent query: Agent query  $Q$  is only sent to the query system when the LLM agent violates given safety constraints in its next action. After a violation is detected by the constraint monitor, safety constraints in NL, valid actions and invalid actions, and their corresponding truth value changes are provided in the prompt together with task specifications, which explicitly instruct the language model to explain the reason for the violation.

The output generated by the query system in response to the agent query,  $l_A$ , will be appended to the original prompt in the form of error messages that the LLM agent uses for decision-making. An example prompt for generating the first error message in the prompt in the previous section could be as follows (task specification and environment information omitted for simplicity):

Constraints: [“you have to enter living room before bathroom”, “you have to enter bedroom before going into living room”]  
 Invalid action: walk to bathroom  
 State change:  
 Safe: !agent\_at(bedroom) & !agent\_at(bathroom) & !agent\_at(livingroom)  
 Violated: !agent\_at(bedroom) & agent\_at(bathroom) & !agent\_at(livingroom)  
 Reason of violation:

#### IV. EXPERIMENTS

The aim of our evaluation is to test the hypothesis that our safety chip approach can reduce the frequency of safety constraint violations without decreasing the frequency of task completion relative to the base models without a safety chip. Experiments are conducted in VirtualHome environment and on the real robot (see section V). We aim to draw comparisons between the proposed method and baseline method where both goal specifications and constraints are fed together to LLM agents.



### A. VirtualHome Environment & Tasks

VirtualHome [39] is a multi-agent platform that supports simulated household tasks in multiple embodied environments. We adopted 20 household tasks from two environments, e.g., “Put salmon into the fridge,” and five satisfiable constraints for each task, e.g., “You must close the fridge if you have ever opened it.” All constraints fall in avoidance and trigger patterns formulated by Menghi et al. [40]. Task difficulty is controlled as each task is accompanied by an in-context example that has the exact ground truth action sequence with only the goal specification paraphrased, and all tasks are ensured to be achievable under the constraints, to draw a fair comparison on the ability to reason about safety constraints between different systems. During the experiment, we access the simulator for information such as locations and states of objects for truth values.

### B. LLM Agents

We proposed the following three LLM agents:

**Base Model** is the foundation of the other two agents. We develop the model based on SayCan [41] and LLM-Planner [27]: in the prompt, we provide environmental information and available actions. During inference, the agent generates one action at a time, grounds it into available actions, and appends it to the prompt for the subsequent generation loop. Though the base model is unaware of any safety constraint, we evaluate its performance as a reference. **Base Model with NL Constraints** (NL Constraints) is instructed through task specification to complete the given task while obeying the safety constraints. It takes as input the safety constraints together with goal specifications. All reasoning over safety constraints and planning happens internally within the LLM.

**Base Model with Safety Chip** (Safety Chip) is the system we developed by plugging the safety chip in the base model. While the base model is unaware of the safety constraints at the beginning, it is instructed to be vigilant for the error messages regarding safety violations and ready to re-generate the next action. At the same time, the safety chip keeps tracking the states of the robot agent, pruning unsafe actions, and conveying the information regarding violation through agent query, as stated in Section III.

Throughout the experiment, GPT4 [42] (gpt-4-0613) is used as the language model for all prompting tasks.

### C. Experimental Setup

We conducted two sets of experiments, *Four Room* and *Mobile Manipulation*, for examining the ability of LLM agents to reason constraints and evaluating the performance under realistic and everyday constraints. For both sets of experiments, we applied an increasing number of constraints from one to five, and we evaluated the performance based on *success rate*, if the end state matches the goal state and is executable by the simulator, and *safety rate*, if ever the agent violates the safety constraints and enters an invalid state specified by the automaton, i.e., a safe execution could be unsuccessful, and vice versa.

For NL constraints and LTL formulas, we use the translation system of safety chip to convert safety constraints specified by LTL experts in natural language to LTL formulas, and then the experts are asked to verify the generated LTL formulas. In this experiment, the experts are also asked to provide accurate NL constraints based on the verified LTL formulas in addition to the original ones for the NL Constraints baseline to assess the effect of ambiguity in NL for a comprehensive evaluation. To be more specific on the two sets of experiments:

**Four Room** is proposed to fairly evaluate the reasoning ability over safety constraints of various systems, where all NL constraints are in similar formats, e.g., “Don’t go to kitchen before bathroom,” and only consist of a navigational predicate,  $\{agent\_at()\}$ , and rooms in  $\{Kitchen, Bathroom, Bedroom, Livingroom\}$  as propositions. Intuitively, we expect the LLM agents to focus more on reasoning over constraints when provided with straightforward NL constraints, rather than on language understanding.

**Mobile Manipulation** tests the system in more realistic scenarios, where we extend the constraints to mobile manipulation domain by incorporating six more predicates,  $\{is\_switchedon(), is\_open(), is\_grabbed(), is\_touched(), is\_on(), is\_in()\}$ , representing the status and relationships of objects, and accepting both objects and locations as arguments. The NL constraints are generally more diverse than in the Four Room experiment, and we also provide a list of available objects, 105 objects per environment on average, in the prompt for the agent to interact with.




Fig. 4: Average Safety Rate by Number of Constraints for VirtualHome

### D. Results

As Table I and Figure 4 show, Safety Chip could achieve 100% safety rate with expert-verified LTL formulas in both experiments, which significantly outperformed the other baselines, especially under a larger number of constraints. Besides, even without expert verification, there is still a large margin over the other baselines, though the success rate is affected because of the mistranslated safety constraints. On the contrary, the NL Constraints baseline struggles to adhere to the safety constraint even with expert-provided NL constraints. Besides, we observed that the success rate of Base Model is lower than the other two agents that have access to the safety constraints, which might indicate the safety constraints could contain useful information that can help decision-making in return.

TABLE I. Average success rate and safety rate of VirtualHome Experiments

	Four Room		Mobile Manipulation	
	Success Rate	Safety Rate	Success Rate	Safety Rate
Base Model	90%	38%	94%	30%
NL Constraints	100%	56%	100%	56%
NL Constraints*	100%	54%	96%	62%
Safety Chip	76%	84%	88%	94%
Safety Chip*	98%	100%	98%	100%

\* expert-verified

## V. ROBOT DEMO

To assess the potential of the proposed system for handling complex safety constraints in practice, we doubled the maximum number of constraints applied to each task and deployed the system on a Spot [43] robot together with two baselines to draw comparisons.

### A. Experimental Setup

We build an indoor environment consisting of 23 objects and landmarks, and the environmental information is stored in the scanned graph from the Spot robot. We define two mobile manipulation tasks with ten safety constraints covering six patterns in avoidance and trigger pattern defined by Menghi et al. [40] per task, and we progressively increase the number of constraints applied to the task. All safety constraints are generated by LTL experts through natural language, and we assume no mistranslation in LTL formulas verified by the experts and no mechanical failure in the experiment, since the goal is to assess the agent's ability to plan under complex constraints.

For baseline comparisons, we adopt the NL Constraints and the Safety Chip agents which are the same as the VirtualHome experiments, and we additionally implement Code as Policies [26] for comparison as a code generation-based LLM agent. We provide predefined mobile manipulation skills, access to object states, and access to the planner of the robot to all baselines, and we build a precondition checker for NL Constraints and Safety Chip to infer physical limitations that indicate available actions at each state, e.g., "pick up mail" is unavailable when the robot agent is not around the mail, while Code as Policies is inherently hard to incorporate closed-loop feedback modules for reprompting due to its code format. During the experiment, each task is run twice, and the temperature of the LLM is set to zero to minimize the effect of stochasticity of the LLM.

### B. Results

As Table II shows, with the assistance of LTL experts, we achieve 100% safety rate without impacting the success rate. When the number of constraints surpasses five, we observe a drastic decrease in the safety rate for the two baselines: NL Constraints achieves 0%, and Code as Policies achieves 40%. Compared with their performance when constraints are fewer than five, where NL Constraints achieves 75% and Code as Policies achieves 70%, we can observe the pattern that the two LLM agents struggle to scale up as the complexity of constraints increases. In addition, three

TABLE II. Average success rate and safety rate of physical robot

	Success Rate	Safety Rate
Code as Policies	53%	55%
NL Constraints	98%	34%
Safety Chip	98%	100%

task abortion tests with safety constraints contradicting the task itself are also conducted to learn the agent's behavior under self-contradicted commands. As a result, we find NL Constraints and Code as Policies still try to accomplish the task even when instructed to obey the safety constraints, which leads to safety violations, while our system can catch the contradicting constraints and abort properly without visiting any unsafe state.

## VI. CONCLUSION

To address the safety concern of deploying LLM agents in practice, we introduced the safety chip that supports customizable constraints by simultaneously encoding NL constraints, conveying the encoded safety constraints to LLM agents, and monitoring and assisting the decision-making process of an LLM agent. From the results, we can draw three conclusions:

- Safety constraints are challenging and distracting for LLM agents to handle without external help, as the complexity of constraints increases.
- Verifiable formal language such as LTL is generally more reliable and interpretable than latent representations in fully end-to-end approaches, particularly concerning the safety aspects of robot agents.
- Ablating things, such as safety constraints, from different abstraction levels holds the potential to help LLM agents concentrate on the main task and effectively function as a "brain".

Lastly, we emphasize the pressing need for more attention to safety aspects from the LLM agents community, as these considerations are indispensable for any practical utilization: when safety becomes a critical issue, the reasoning skills of LLMs or other neural network-based models alone cannot guarantee to satisfy safety standards, and hence, we proposed a potential solution based on formal verification. On a higher level, the proposed "safety chip" also has the potential to assist other frameworks beyond the LLM agent domain in achieving safety standards, and we hope to provide a new perspective on viewing and handling safety constraints for various robotic tasks.

## VII. ACKNOWLEDGEMENT

The authors would like to thank Eric Rosen and Kaiyu Zheng for their thoughtful suggestions. This work is supported by the Air Force Office of Scientific Research under award number FA9550-21-1-0214 and the Office of Naval Research under award number N00014-22-1-2592.

# REFERENCES

- [1] L. Weng, "Llm-powered autonomous agents," *lilianweng.github.io*, Jun 2023. [Online]. Available: <https://lilianweng.github.io/posts/2023-06-23-agent/>
- [2] L. Wang, C. Ma, X. Feng, Z. Zhang, H. Yang, J. Zhang, Z. Chen, J. Tang, X. Chen, Y. Lin, W. X. Zhao, Z. Wei, and J.-R. Wen, "A survey on large language model based autonomous agents," 2023.
- [3] International Organization for Standardization/International Electrotechnical Commission, "Functional safety of electrical/electronic/programmable electronic safety-related systems - part 1: General requirements," 2010.
- [4] S. Gu, L. Yang, Y. Du, G. Chen, F. Walter, J. Wang, Y. Yang, and A. Knoll, "A review of safe reinforcement learning: Methods, theory and applications," 2023.
- [5] B. Greenman, S. Saarinen, T. Nelson, and S. Krishnamurthi, "Little tricky logic: Misconceptions in the understanding of ltl," *The Art, Science, and Engineering of Programming*, vol. 7, 2023.
- [6] A. Pnueli, "The temporal logic of programs," in *18th Annual Symposium on Foundations of Computer Science (sfcs 1977)*. IEEE, 1977, pp. 46–57.
- [7] A. Shah, S. Li, and J. Shah, "Planning with uncertain specifications (puns)," *IEEE Robotics and Automation Letters*, vol. 5, no. 2, pp. 3414–3421, 2020.
- [8] J. X. Liu, A. Shah, E. Rosen, G. Konidaris, and S. Tellex, "Skill transfer for temporally-extended task specifications," *arXiv preprint arXiv:2206.05096*, 2022.
- [9] A. Pacheck and H. Kress-Gazit, "Physically-feasible repair of reactive, linear temporal logic-based, high-level tasks," 2022.
- [10] H. Kress-Gazit, M. Lahijanlian, and V. Raman, "Synthesis for robots: Guarantees and feedback for robot behavior," *Annu. Rev. Control. Robotics Auton. Syst.*, vol. 1, pp. 211–236, 2018. [Online]. Available: <https://api.semanticscholar.org/CorpusID:65105626>
- [11] P. Thati and G. Roşu, "Monitoring algorithms for metric temporal logic specifications," *Electronic Notes in Theoretical Computer Science*, vol. 113, pp. 145–162, 2005.
- [12] M. Foughali, S. Bensalem, J. Combaz, and F. Ingrand, "Runtime verification of timed properties in autonomous robots," in *2020 18th ACM-IEEE International Conference on Formal Methods and Models for System Design (MEMOCODE)*. IEEE, 2020, pp. 1–12.
- [13] T. Reinbacher, K. Y. Rozier, and J. Schumann, "Temporal-logic based runtime observer pairs for system health management of real-time systems," in *Tools and Algorithms for the Construction and Analysis of Systems: 20th International Conference, TACAS 2014*. Springer, 2014, pp. 357–372.
- [14] H. Barringer, A. Goldberg, K. Havelund, and K. Sen, "Rule-based runtime verification," in *Verification, Model Checking, and Abstract Interpretation: 5th International Conference, VMCAI 2004 Venice, Italy, January 11-13, 2004 Proceedings 5*. Springer, 2004, pp. 44–57.
- [15] P. Moosbrugger, K. Y. Rozier, and J. Schumann, "R2u2: monitoring and diagnosis of security threats for unmanned aerial systems," *Formal Methods in System Design*, vol. 51, pp. 31–61, 2017.
- [16] H. Kress-Gazit, K. Eder, G. Hoffman, H. Admoni, B. Argall, R. Ehlers, C. Heckman, N. Jansen, R. Knepper, J. Křetínský *et al.*, "Formalizing and guaranteeing human-robot interaction," *Communications of the ACM*, vol. 64, no. 9, pp. 78–84, 2021.
- [17] X. Mao, X. Li, Y. Huang, J. Shi, and Y. Zhang, "Programmable logic controllers past linear temporal logic for monitoring applications in industrial control systems," *IEEE Transactions on Industrial Informatics*, vol. 18, no. 7, pp. 4393–4405, 2022.
- [18] P. A. Lasota, T. Fong, and J. A. Shah, "A survey of methods for safe human-robot interaction," *Foundations and Trends® in Robotics*, vol. 5, no. 4, pp. 261–349, 2017. [Online]. Available: <http://dx.doi.org/10.1561/23000000052>
- [19] M. Ahn, A. Brohan, N. Brown, Y. Chebotar, O. Cortes, B. David, C. Finn, C. Fu, K. Gopalakrishnan, K. Hausman, A. Herzog, D. Ho, J. Hsu, J. Ibarz, B. Ichter, A. Irpan, E. Jang, R. J. Ruano, K. Jeffrey, S. Jesmonth, N. J. Joshi, R. Julian, D. Kalashnikov, Y. Kuang, K.-H. Lee, S. Levine, Y. Lu, L. Luu, C. Parada, P. Pastor, J. Quiambao, K. Rao, J. Rettinghouse, D. Reyes, P. Sermanet, N. Sievers, C. Tan, A. Toshev, V. Vanhoucke, F. Xia, T. Xiao, P. Xu, S. Xu, M. Yan, and A. Zeng, "Do as i can, not as i say: Grounding language in robotic affordances," 2022.
- [20] W. Huang, F. Xia, T. Xiao, H. Chan, J. Liang, P. Florence, A. Zeng, J. Tompson, I. Mordatch, Y. Chebotar, P. Sermanet, N. Brown, T. Jackson, L. Luu, S. Levine, K. Hausman, and B. Ichter, "Inner monologue: Embodied reasoning through planning with language models," 2022.
- [21] W. Huang, P. Abbeel, D. Pathak, and I. Mordatch, "Language models as zero-shot planners: Extracting actionable knowledge for embodied agents," 2022.
- [22] B. Chen, F. Xia, B. Ichter, K. Rao, K. Gopalakrishnan, M. S. Ryoo, A. Stone, and D. Kappler, "Open-vocabulary queryable scene representations for real world planning," 2022.
- [23] D. Shah, B. Osinski, B. Ichter, and S. Levine, "Lmnav: Robotic navigation with large pre-trained models of language, vision, and action," 2022.
- [24] W. Huang, F. Xia, D. Shah, D. Driess, A. Zeng, Y. Lu, P. Florence, I. Mordatch, S. Levine, K. Hausman, and B. Ichter, "Grounded decoding: Guiding text generation with grounded models for robot control," 2023.
- [25] I. Singh, V. Blukis, A. Mousavian, A. Goyal, D. Xu, J. Tremblay, D. Fox, J. Thomason, and A. Garg, "Prog-prompt: Generating situated robot task plans using large language models," *arXiv preprint arXiv:2209.11302*, 2022.

- [26] J. Liang, W. Huang, F. Xia, P. Xu, K. Hausman, B. Ichter, P. Florence, and A. Zeng, "Code as policies: Language model programs for embodied control," in *arXiv preprint arXiv:2209.07753*, 2022.
- [27] B. Liu, Y. Jiang, X. Zhang, Q. Liu, S. Zhang, J. Biswas, and P. Stone, "Llm+p: Empowering large language models with optimal planning proficiency," 2023.
- [28] Y. Xie, C. Yu, T. Zhu, J. Bai, Z. Gong, and H. Soh, "Translating natural language to planning goals with large-language models," 2023.
- [29] N. Gopalan, D. Arumugam, L. Wong, and S. Tellex, "Sequence-to-Sequence Language Grounding of Non-Markovian Task Specifications," in *Proceedings of Robotics: Science and Systems*, Pittsburgh, Pennsylvania, 2018.
- [30] R. Patel, E. Pavlick, and S. Tellex, "Grounding language to non-markovian tasks with no supervision of task specifications," in *Robotics: Science and Systems*, 2020.
- [31] Y. Chen, R. Gandhi, Y. Zhang, and C. Fan, "Nl2tl: Transforming natural languages to temporal logics using large language models," *arXiv preprint arXiv:2305.07766*, 2023.
- [32] J. Pan, G. Chou, and D. Berenson, "Data-efficient learning of natural language to linear temporal logic translators for robot task specification," *arXiv preprint arXiv:2303.08006*, 2023.
- [33] J. X. Liu, Z. Yang, I. Idrees, S. Liang, B. Schornstein, S. Tellex, and A. Shah, "Lang2tl: Translating natural language commands to temporal robot task specification," 2023.
- [34] H. Cherukuri, A. Ferrari, and P. Spoletini, "Towards explainable formal methods: From ltl to natural language with neural machine translation," in *Requirements Engineering: Foundation for Software Quality*, V. Gervasi and A. Vogelsang, Eds. Cham: Springer International Publishing, 2022, pp. 79–86.
- [35] V. Raman, C. Lignos, C. Finucane, K. C. Lee, M. P. Marcus, and H. Kress-Gazit, "Sorry dave, i'm afraid i can't do that: Explaining unachievable robot tasks using natural language," in *Robotics: science and systems*, vol. 2, no. 1. Citeseer, 2013, pp. 2–1.
- [36] S. S. Raman, V. Cohen, E. Rosen, I. Idrees, D. Paulius, and S. Tellex, "Planning with large language models via corrective re-prompting," *arXiv preprint arXiv:2211.09935*, 2022.
- [37] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao, "React: Synergizing reasoning and acting in language models," 2023.
- [38] Z. Ji, N. Lee, R. Frieske, T. Yu, D. Su, Y. Xu, E. Ishii, Y. J. Bang, A. Madotto, and P. Fung, "Survey of hallucination in natural language generation," *ACM Computing Surveys*, vol. 55, no. 12, pp. 1–38, mar 2023. [Online]. Available: <https://doi.org/10.1145%2F3571730>
- [39] X. Puig, K. Ra, M. Boben, J. Li, T. Wang, S. Fidler, and A. Torralba, "Virtualhome: Simulating household activities via programs," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 2018, pp. 8494–8502.
- [40] C. Menghi, C. Tsigkanos, P. Pelliccione, C. Ghezzi, and T. Berger, "Specification patterns for robotic missions," *IEEE Transactions on Software Engineering*, vol. 47, no. 10, pp. 2208–2224, oct 2021.
- [41] M. Ahn, A. Brohan, N. Brown, Y. Chebotar, O. Cortes, B. David, C. Finn, K. Gopalakrishnan, K. Hausman, A. Herzog *et al.*, "Do as i can, not as i say: Grounding language in robotic affordances," *arXiv preprint arXiv:2204.01691*, 2022.
- [42] OpenAI, "Gpt-4 technical report," 2023.
- [43] Boston Dynamics, "Spot - the agile mobile robot," <https://www.bostondynamics.com/products/spot>.